

Alexandre Kirchmeyer

+1 (609) 608-8200 | akirchme@cs.cmu.edu | [website](#) | [linkedin](#) | [github](#) | [bitbucket](#)

EDUCATION

Carnegie Mellon University – Pittsburgh, PA

M.S. in Machine Learning

École Polytechnique – Palaiseau, France

M.S. in Mathematics and Computer Science

Sep. 2022 – Dec. 2023

QPA: 4.06/4.0

Aug. 2019 – Aug. 2022

GPA: 3.90/4.0

EXPERIENCE

Research Assistant – advised by [Prof. Jia Deng](#), Princeton University

Mar. 2022 – Aug. 2022

Researched the use of a novel oriented 1D convolution operator, to improve the long-range scaling of CNNs at linear cost.

- Implemented personalized CUDA kernels from scratch which are up to 1.5x faster than PyTorch 1D and 30x faster than 2D.
- Trained first state-of-the-art fully 1D ConvNeXt model, demonstrating that our 1D approach can match 2D on image classification, semantic segmentation and object detection. 1st author publication at ICCV 2023. [[paper](#)][[poster](#)][[code](#)]

Research Assistant – advised by [Prof. Deepak Pathak](#), Carnegie Mellon University

Dec. 2022 – Present

Investigated the use of diffusion models for zero-shot 3D reconstruction from one image using a DreamFusion approach.

- Created 100k multiview synthetic dataset by generating partial point clouds with StableDiffusion/MiDaS and doing inpainting.
- Fine-tuned StableDiffusion on synthetic data with LoRA/ControlNet/InstructPix2Pix conditioned generation. [[report](#)][[code](#)]

Exploring sim2real transfer and perception for the task of catching dynamic objects in the context of embodied AI.

- Training RL catching policy in simulation (80% success rate) and transferring to real-world via policy distillation.
- Perception using 3D vision (point cloud, depth, SDF), supervised by dexterous grasping (DexGraspNet losses, eigengrasps).

Research Assistant – advised by [Prof. Andrei Risteki](#), Carnegie Mellon University

Jul. 2023 – Present

Analyzing the trade-off between masking and downstream performance in language models using statistical tools.

Research Assistant – advised by [Pr. Maks Ovsjanikov](#), INRIA, Ecole Polytechnique

Sep. 2021 – Mar. 2022

Examined the use of hyper-networks and deformed implicit fields for the task of non-rigid 3D shape correspondence.

- Re-implemented Microsoft DIF-Net on JAX and evaluated shape correspondence performance on KeypointNet dataset.
- Devised hybrid correspondence approach, halving error threshold to achieve 50% keypoint accuracy. [[report](#)][[code](#)]

Research and Development Intern – [HarfangLab](#) Cybersecurity Startup

Jun. 2021 – Sep. 2021

Implemented HarfangLab's first intrusive antivirus method to detect malicious behavior using DLL injection and API hooking.

- Designed a C++ Windows driver from scratch and automated testing of 10+ malicious behaviors using virtual machines.

PROJECTS

Carnegie Mellon University

Sep. 2022 – Dec. 2022

Introduced concepts to improve model robustness on visual question answering – advised by [Prof. Louis-Philippe Morency](#)

- Proposed to use concepts (color, ...) to improve downstream performance in a teacher-student setup: first train concept teachers with a multimodal knowledge base, then train students with pseudo-labels from teachers and multi-task objective.
- Mined unsupervised multimodal concepts with prefix tuning & contrastive learning, and analyzed explainability with clusters.
- Introduced diffusion-based synthetic causal interventions as data augmentations on the Visual7W dataset. Measured the sim2real and causal gaps and showed that causal interventions improve robustness on the ViLT model by 10%.
- Selected as one of the two *Best Midterm Presentation* for the course out of 24 groups. [[report](#)][[poster](#)][[code](#)]

Generated multi-agent trajectories with diffusion models to improve self-driving multi-agent trajectory planning policies.

- Trained a diffusion and VQ-diffusion model on the task of multi-agent trajectory generation and experimented with different samplers. Experimentally observed that continuous diffusion and deterministic samplers generated better trajectories.

Explored the use of zero-shot vision foundation models for determining driving intent for self-driving applications.

- Fine-tuned a YOLOv8 object detector and evaluated zero-shot vision foundation models like DETIC, Grounding DINO and ODISE on the task of identifying rear vehicle signals.

SKILLS

Programming Languages: C, C++, Python, CUDA, JavaScript, OCaml, Lua, Java, OpenGL, OpenCL, Vulkan, SQL

Software: PyTorch, JAX, ArchLinux, ROS, ReactJS, Slurm, Git, LXC, WandB, transformers, diffusers, IsaacGym, Cutlass

RELEVANT COURSEWORK

Machine Learning (PhD)

Competitive Programming

Computer Vision (PhD)

General Algebra & Galois Theory

Natural Language Processing

Convex Optimization (PhD)

3D Computer Vision

Real Analysis & Variational Methods

Reinforcement Learning (PhD)

Probabilistic Graphical Models (PhD)

Multimodal ML (PhD)

Advanced Quantum Physics

AWARDS

- [Ranked](#) 4th out of 107 at the ACM ICPC SWERC 2020-2021 European Algorithmics Contest Mar. 2021
- [Ranked](#) 3rd out of 12 at the ICPC Ecole Polytechnique Championship 2020 Jun. 2020
- [Ranked](#) 8th out of 62 at the French-Australian Regional Informatics Olympiad 2017 Mar. 2017
- [USACO Platinum](#) level Dec. 2016

PUBLICATIONS

- Alexandre Kirchmeyer, Jia Deng, Convolutional Networks with Oriented 1D Kernels**, ICCV, 2023. [[paper](#)] [[code](#)]